



AWS Auto Scaling - Complete Guide

From Single Instance to Highly Available Scalable Infrastructure



The Problem: Single Instance Limitation

Before Auto Scaling

Current Setup:

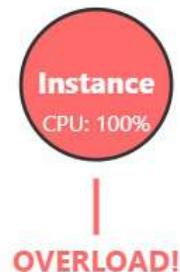
- 1 EC2 Instance running
- Users increase → CPU 100%
- Application slows down
- Users get timeouts
- Manual scaling needed
- No redundancy

After Auto Scaling

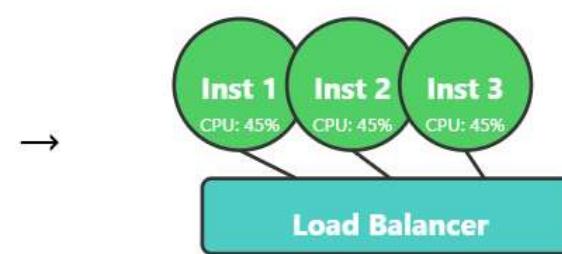
Desired Setup:

- Multiple EC2 Instances
- Auto scales based on load
- Always optimal performance
- No user interruption
- Automatic scaling policy
- High availability

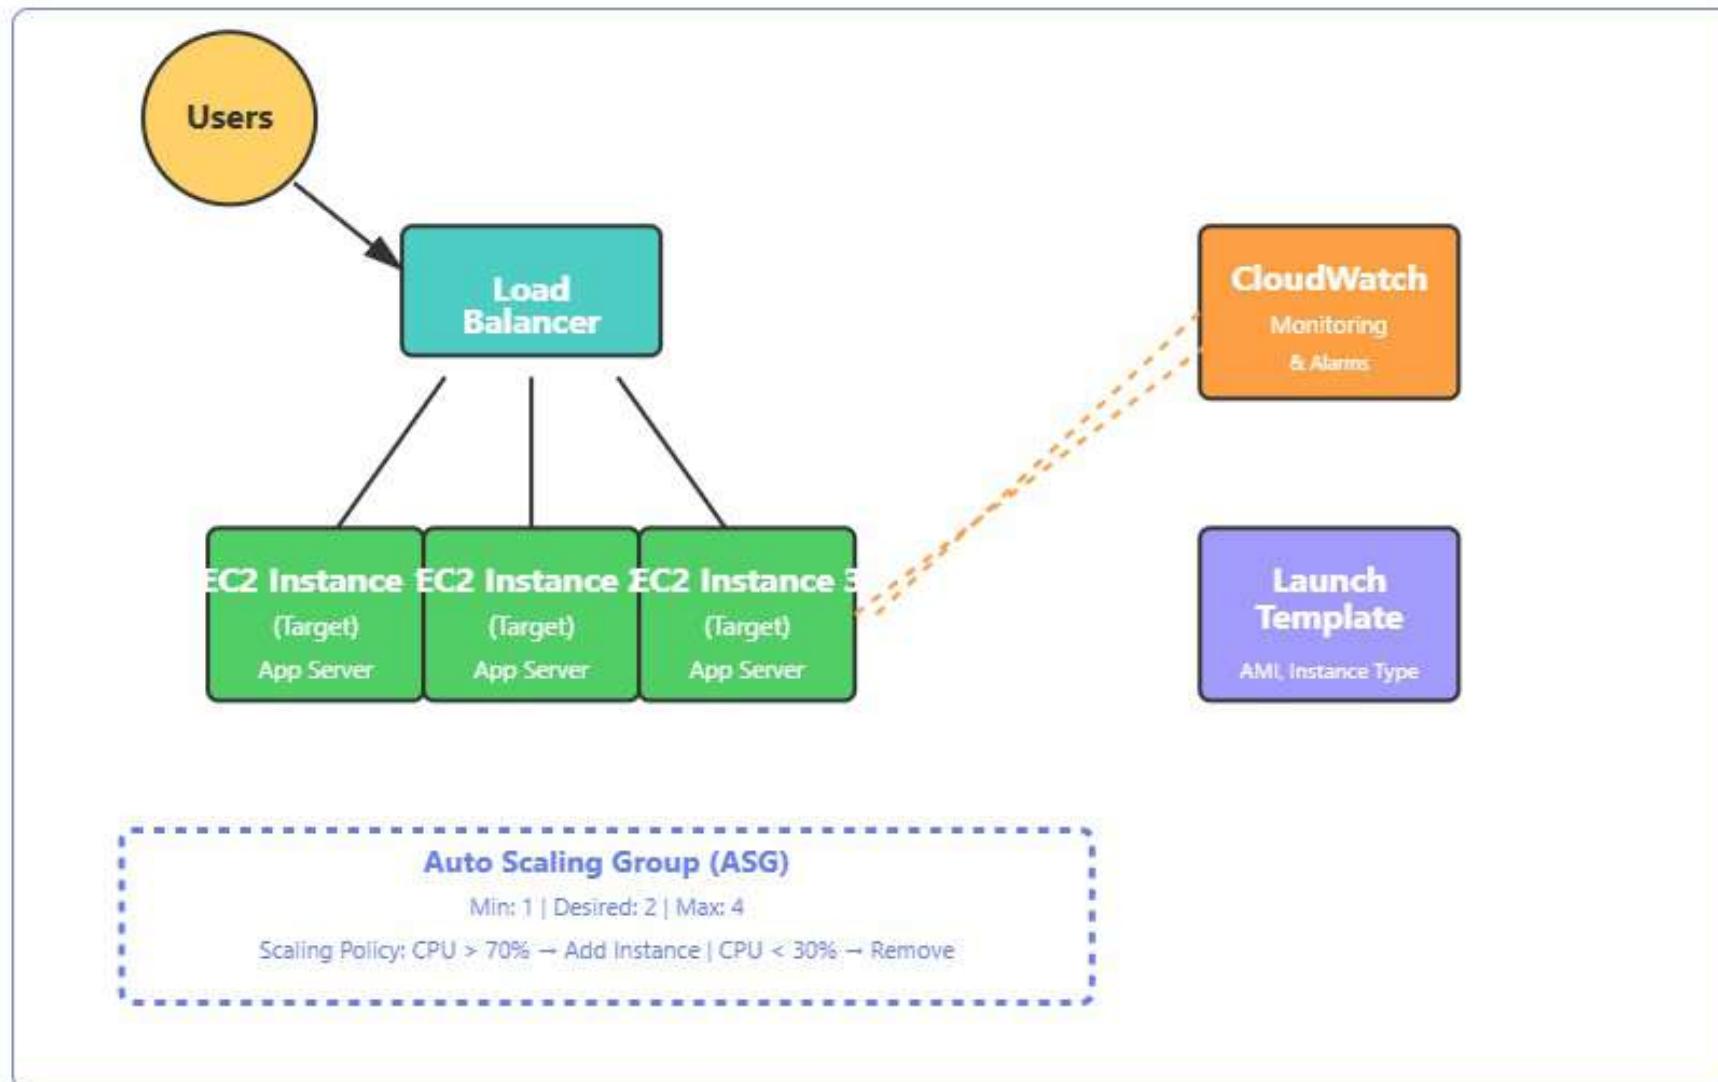
Before Auto Scaling



After Auto Scaling



Auto Scaling Architecture





Complete Step-by-Step Setup Guide

Step 1: Create a Launch Template

1 Navigate to Launch Templates in EC2 Console

```
AWS Console → EC2 → Launch Templates → Create Launch Template
```

Configuration:

Name: my-app-template

AMI: Select your base image (Ubuntu 22.04 LTS recommended)

Instance Type: t3.micro (or t3.small for production)

Key Pair: Select your existing key pair

Security Group: Create/select one allowing HTTP (80), HTTPS (443), SSH (22)

Step 2: Create Auto Scaling Group (ASG)

2 Create Auto Scaling Group

```
AWS Console → EC2 → Auto Scaling → Auto Scaling Groups → Create
```

Configuration:

Name: my-app-asg

Launch Template: Select my-app-template

VPC: Select your VPC

Subnets: Select 2-3 subnets (different AZs for HA)

⌚ Key ASG Settings:

Min Capacity: 1 (minimum instances always running)

Desired Capacity: 2 (initial number of instances)

Max Capacity: 4 (maximum during peak load)

Step 3: Configure Load Balancer

3 Create Application Load Balancer (ALB)

```
AWS Console → EC2 → Load Balancing → Load Balancers → Create
```

Configuration:

Type: **Application Load Balancer**

Name: my-app-alb

Scheme: **Internet-facing**

Subnets: Select same subnets as ASG

Listener Port: **80 (HTTP)**

Create Target Group:

Name: my-app-targets

Target Type: **Instances**

Protocol: **HTTP**

Port: **80**

Health Check Path: /

Healthy Threshold: 2

Unhealthy Threshold: 3

Interval: 30 seconds

Step 4: Attach Load Balancer to ASG

4 Link ALB to Auto Scaling Group

```
Edit ASG → Load Balancing → Select Target Group
```

This connects:

Incoming traffic through ALB

Distributed to ASG instances

Health checks through target group

Step 5: Create Scaling Policies

5 Configure Scaling Policies (Target Tracking)

Edit ASG → Automatic Scaling → Target Tracking Scaling Policies

Policy 1: Scale Out (Add Instances)

Metric Type: **Average CPU Utilization**

Target Value: **70%**

When CPU > 70%, ASG adds instances

Policy 2: Scale In (Remove Instances)

Metric Type: **Average CPU Utilization**

Target Value: **30%**

When CPU < 30%, ASG removes instances

⌚ Important Cooldown Settings:

Scale-out Cooldown: 60 seconds (wait before adding more)

Scale-in Cooldown: 300 seconds (wait before removing)

Step 6: Monitor with CloudWatch

6 Setup CloudWatch Monitoring

AWS Console → CloudWatch → Dashboards → Create Custom Dashboard

Metrics to Monitor:

ASG GroupInServiceInstances

ASG GroupTerminatingInstances

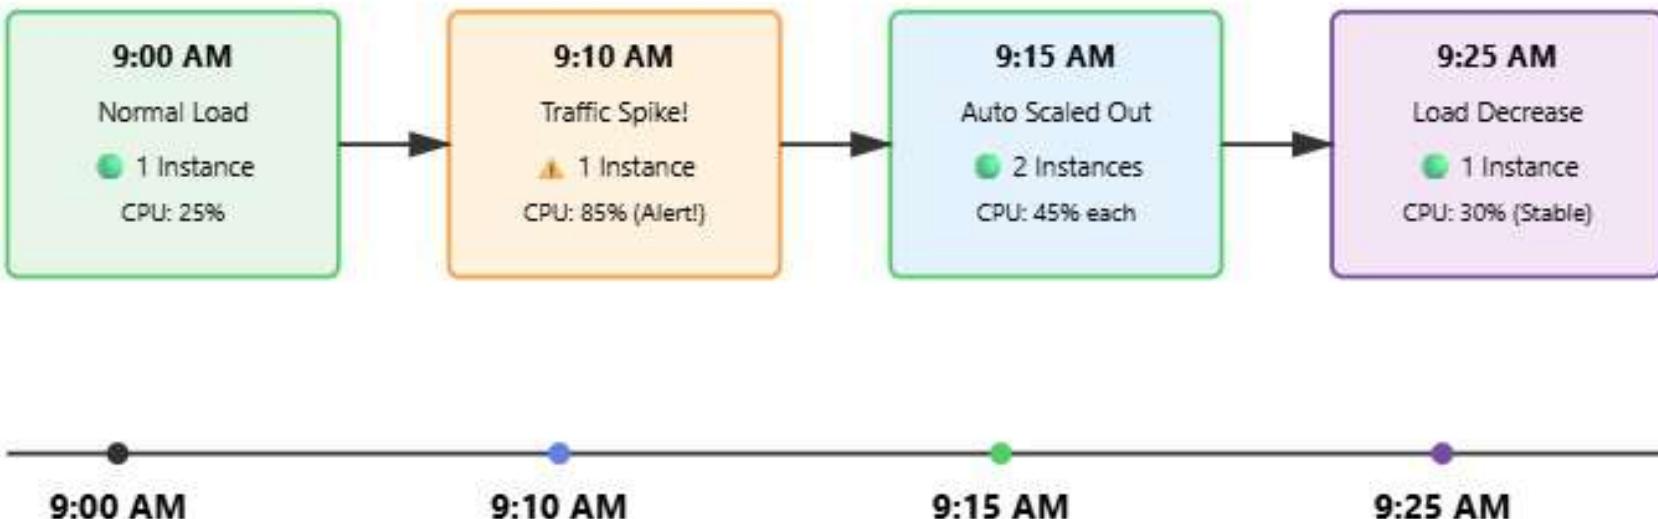
Average CPU Utilization

ALB TargetResponseTime

ALB RequestCount



How Auto Scaling Works - Timeline





Practical Example - Setup from Scratch

Your Scenario: E-commerce Website

You have a website currently running on 1 t3.micro instance. During sales, traffic can spike 10x. You need auto scaling.

Solution Implementation:

| Component | Configuration | Reason |
|---------------------|----------------------------------|--------------------------------------|
| Instance Type | t3.micro (free tier) or t3.small | Good for web apps, cost-effective |
| Min Capacity | 1 | Always have 1 running (cost savings) |
| Desired Capacity | 2 | Start with 2 for redundancy |
| Max Capacity | 5 | Cap at 5 to control costs |
| Scale Out Threshold | CPU \geq 70% | Add instance when busy |
| Scale In Threshold | CPU \leq 30% | Remove instance when idle |

Implementation Script (AWS CLI)

```
# 1. Create Launch Template aws ec2 create-launch-template \
--launch-template-name my-ecommerce-template \
--version-description "Web app template" \
--launch-template-data '{ "ImageId": "ami-0c55b159cbfafef0",
"InstanceType": "t3.micro", "KeyName": "my-key-pair", "SecurityGroupIds": ["sg-12345678"] }' # 2. Create Auto Scaling Group aws autoscaling create-auto-scaling-group \
--auto-scaling-group-name my-ecommerce-asg \
--launch-template LaunchTemplateName=my-ecommerce-template \
--min-size 1 \
--desired-capacity 2 \
--max-size 5 \
--availability-zones us-east-1a us-east-1b \
--target-
```