

# Mock Interview Guide

## Terraform

### Instructions for Interviewer:

- You are playing the role of **interviewer**. Use this guide as a script.
  - Ask each question one at a time. Follow the steps: **Definition → Details → Scenario → Follow-up**.
  - If the interviewee struggles, use the **hint**.
  - The goal is to keep it conversational and practical. Help the interviewee think and express their learning.
  - **colors assigned:** Questions Answers Hint
- 

### Freshers - Level

### Terraform

### (10 Easy DevOps Interview Questions)

---

#### 1. “What is Terraform used for?”

✓ Expected Answer: “Terraform is an open-source tool used to automate the provisioning of cloud infrastructure using code.”

Hint: “*Think of Terraform like writing blueprints for your cloud. Instead of clicking around in AWS or Azure, you write code that builds the setup for you.*”

#### 2. “What does IaC mean in Terraform?”

✓ Expected Answer: “IaC stands for Infrastructure as Code. It means managing and provisioning infrastructure through code instead of manually.”

**Hint:** “Imagine describing your servers and networks the same way you write code for apps — then letting a tool set them up for you.”

### 3. “What file format does Terraform use?”

✓ Expected Answer: “Terraform uses files written in HashiCorp Configuration Language (HCL) with .tf extension.”

**Hint:** “You’re writing configuration, not code logic. It’s stored in plain-text files with a specific extension.”

### 4. “What is terraform init?”

✓ Expected Answer: “It initializes a Terraform working directory and downloads required providers.”

**Hint:** “The very first step before you apply any changes — it sets up your project folder and downloads cloud plugins.”

### 5. “What does terraform plan do?”

✓ Expected Answer: “It shows what changes Terraform will make to the infrastructure before applying them.”

**Hint:** “Want to see what will happen before you actually build something? This command gives you that preview.”

### 6. “How do you apply the Terraform configuration?”

✓ Expected Answer: “By using the terraform apply command.”

**Hint:** “Once you’ve planned your infrastructure, this command actually builds or updates it.”

## 7. “How do you destroy resources created by Terraform?”

❖ Expected Answer: “By running the `terraform destroy` command.”

**Hint:** “Used when you want to clean up everything you’ve provisioned using your `.tf` files.”

## 8. “What is a Terraform provider?”

❖ Expected Answer: “A provider is a plugin that Terraform uses to interact with APIs of cloud platforms like AWS, Azure, GCP.”

**Hint:** “This is how Terraform knows how to talk to AWS or Azure — think of it as the bridge between Terraform and the cloud.”

## 9. “Where are Terraform states stored?”

❖ Expected Answer: “By default, locally in a file called `terraform.tfstate`, but it can be stored remotely.”

**Hint:** “This file keeps track of what’s already built, so Terraform knows what needs changing.”

## 10. “What is a variable in Terraform?”

❖ Expected Answer: “A variable is a reusable value defined in configuration that can be customized.”

---

Hint: “*Think of it like a placeholder — instead of hardcoding values like AMI ID or region, you make it flexible.*”

---

## SCENARIO-BASED INTERVIEW QUESTIONS

---

**1. Ask: “*You’ve run terraform plan, but nothing changed in your cloud environment. What do you think might have been missed?***

✓ Expected: *The candidate likely forgot to run terraform apply after planning.*

Hint: *What command actually makes the changes?*

**2. Ask: “*Suppose you updated a value in terraform.tfvars, but running terraform plan shows no difference. What could be the reason?***

✓ Expected: *Maybe the variable isn’t being referenced in the code, or terraform plan wasn’t re-run properly.*

Hint: *Did the variable actually connect to anything in your config?*

**3. Ask: “*You changed the backend configuration, but Terraform throws an error during init. Why does this happen?***

✓ Expected: *Backend config is locked after first init — you must reinitialize with terraform init -reconfigure.*

Hint: *What’s the safe way to tell Terraform you changed how it stores state?*

**4. Ask: “Imagine you’ve cloned someone else’s Terraform project and terraform plan fails immediately. What’s your first step?”**

✓ Expected: *Run terraform init to initialize providers and setup.*

Hint: *What command prepares the environment for use?*

**5. Ask: “Your Terraform script tries to delete an S3 bucket but fails. What could be the reason?”**

✓ Expected: *Terraform can’t delete an S3 bucket unless it’s empty. Manual cleanup or force-delete setup is needed.*

Hint: *Did the bucket contain files during deletion?*

---

## PROJECT-BASED INTERVIEW QUESTIONS

---

**1. Ask: “If I ask you to provision a basic EC2 instance using Terraform, how would you go about it?”**

✓ Expected: *Use aws\_instance with AMI ID and instance type in a .tf file, then run init, plan, and apply.*

Hint: *Think: one resource, three core commands.*

**2. Ask: “I want you to create a simple S3 bucket using Terraform. What would your workflow look like?”**

✓ Expected: *Define aws\_s3\_bucket, provide a unique name and region, then apply the plan.*

Hint: *What's the minimum info needed to make a bucket?*

3. Ask: *“Let’s say we need to allow HTTP and SSH traffic to a server. How would you create the security group for it in Terraform?”*

✓ Expected: *Use aws\_security\_group with ingress rules for ports 22 and 80.*

Hint: *Which block helps define allowed connections?*

4. Ask: *“How would you structure your Terraform files for a small, single-module project?”*

✓ Expected: *Use main.tf for resources, variables.tf, and outputs.tf. Use terraform.tfvars for values.*

Hint: *How do you make your code readable and organized from day one?*

---

# **Medium - Level**

## **Terraform**

### **(DevOps Interview Questions - 1 to 2 Years Experience)**

---

#### **1. “What is the difference between terraform apply and terraform plan?”**

✓ Answer:

**terraform plan shows the changes Terraform will make without applying them.**

**terraform apply actually performs those changes on the infrastructure.**

Hint: *One is a dry run, the other makes it real.*

#### **2. “What is the purpose of terraform state command?”**

✓ Answer:

**It lets you inspect, move, or remove items in the current Terraform state.**

**Useful for advanced debugging and resource tracking.**

Hint: *You're looking into or fixing the .tfstate file.*

#### **3. “How do you manage multiple environments like dev, stage, and prod?”**

✓ Answer:

**Use Terraform workspaces, or separate directories/modules with**

**different variable files.  
This ensures isolation and cleaner configuration.**

**Hint:** *Think reuse, but with separation.*

#### **4. “What are Terraform modules and why are they useful?”**

**✓ Answer:**  
**Modules are reusable components that help organize and simplify complex infrastructure.**  
**They reduce duplication and enforce consistency.**

**Hint:** *Think of them like reusable functions in code.*

#### **5. “How do you manage secrets securely in Terraform?”**

**✓ Answer:**  
**Avoid hardcoding secrets; use environment variables, remote backends, or tools like Vault.**  
**Sensitive values should never be stored in code.**

**Hint:** *Security 101: don’t put passwords in .tf files.*

#### **6. “What are data sources in Terraform?”**

**✓ Answer:**  
**Data sources let Terraform fetch existing info from your provider (e.g., AMI IDs).**  
**They’re read-only and don’t create resources.**

**Hint:** *You want to use data that already exists.*

## 7. “Difference between count and for\_each?”

✓ Answer:

count is for simple replication using integers.

for\_each gives more control using maps or sets with unique keys.

Hint: *Want named resources tied to specific data?*

## 8. “How does remote state work?”

✓ Answer:

Remote state stores .tfstate files in services like S3 or Terraform Cloud.

It enables team collaboration and locking.

Hint: *You need shared state across a team.*

## 9. “How to handle resource dependencies in Terraform?”

✓ Answer:

Terraform automatically detects dependencies using references.

You can manually specify them using depends\_on.

Hint: *Sometimes you need to control execution order.*

## 10. “What is the lifecycle block in Terraform?”

✓ Answer:

It lets you customize resource behavior, like forcing recreation or preventing deletion.

Useful for zero-downtime updates.

Hint: *Control how resources are destroyed or replaced.*

---

## SCENARIO-BASED INTERVIEW QUESTIONS

---

**1. Ask: “You’re working in a team and both you and a teammate tried applying changes simultaneously. One apply failed with a lock error. Why?”**

❖ Expected: *Terraform state was locked. Only one apply can modify the state at a time.*

Hint: *Think about how Terraform avoids conflicting writes.*

**2. Ask: “You’re using a remote backend with S3 and DynamoDB for state locking. What happens if someone deletes the DynamoDB table?”**

❖ Expected: *State locking would fail, risking concurrent updates to the state.*

Hint: *What part protects your state from being changed at the same time?*

**3. Ask: “You used count to create multiple instances, but now want to delete one without affecting others. What challenge will you face?”**

❖ Expected: *count relies on index, so removing one causes shifting. Use for\_each for better resource targeting.*

Hint: *Count is index-based. What’s better for selective removal?*

**4. Ask: “You changed a module but the plan didn’t show any updates. What might be the reason?”**

✓ Expected: *You may have forgotten to run terraform get -update to refresh the module.*

Hint: *What command ensures module changes are recognized?*

**5. Ask: “You’re using terraform import to bring an existing AWS resource under Terraform control. What’s the limitation here?”**

✓ Expected: *It only imports state — no configuration. You must manually write the matching .tf code.*

Hint: *Import pulls in state, not code.*

---

## PROJECT-BASED INTERVIEW QUESTIONS

---

**1. Ask: “You’ve been asked to create a reusable Terraform module for provisioning EC2 instances. How would you structure it?”**

✓ Expected: *Use main.tf, variables.tf, and outputs.tf inside a separate module folder. Then call it using module block in root.*

Hint: *Think modular: input → reusable logic → output.*

**2. Ask:** “*How would you design a Terraform workflow that deploys both a VPC and EC2 instance — but only if the VPC was created successfully?*”

✓ Expected: *Define the EC2 instance with a dependency on VPC outputs. Terraform builds the dependency tree automatically.*

Hint: *Let Terraform handle the dependency chain using references.*

**3. Ask:** “*You’re provisioning infra for two environments: dev and prod. How would you structure the Terraform project to support both?*”

✓ Expected: *Use workspaces or maintain separate folders for dev/ and prod/ with shared modules.*

Hint: *Environments need separation — by workspace or folder.*

**4. Ask:** “*Suppose you need to deploy a public and a private subnet using Terraform. What AWS resource will help you route traffic properly?*”

✓ Expected: *Use route tables: one public (with internet gateway), one private (with NAT gateway if needed).*

Hint: *Routing rules are key to network segregation.*

---

# **Hard - Level**

## **Terraform**

### **(DevOps Interview Questions - 3+ Years Experience)**

---

#### **1. “How do you handle state locking in Terraform?”**

✓ **Answer:**

**State locking prevents concurrent modifications to the state file.**  
**It's automatically handled in backends like S3 with DynamoDB or Terraform Cloud.**

*Hint: What if two people run apply at the same time?*

#### **2. “What are the benefits of using Terraform Cloud or Terraform Enterprise over local execution?”**

✓ **Answer:**

**They provide remote state storage, policy enforcement, RBAC, collaboration, and cost estimation.**  
**Ideal for teams working at scale.**

*Hint: Think collaboration, security, governance, and audit trails.*

#### **3. “How do you manage secrets in Terraform when using S3 backend with version control?”**

✓ **Answer:**

**Avoid committing secrets; use sensitive = true, environment variables, or reference secrets from external tools like Vault or SSM.**  
**Never store secrets in .tf files or state.**

*Hint: Even state files can leak secrets — what's your strategy?*

#### **4. “What is Terraform drift detection and how do you handle it?”**

✓ Answer:

**Drift detection is identifying changes in infrastructure that were made outside Terraform.**

**You detect it using terraform plan and resolve by reconciling or reapplying.**

*Hint: What if someone manually modified a resource in AWS?*

#### **5. “How can you modularize Terraform code for large-scale environments?”**

✓ Answer:

**Use well-structured modules with clear inputs/outputs, follow DRY principles, and organize reusable logic per resource/domain.**

**Use terragrunt or mono-repos for additional control.**

*Hint: You're managing 200+ resources — how do you stay sane?*

#### **6. “What are dynamic blocks and when should you use them?”**

✓ Answer:

**Dynamic blocks generate nested blocks programmatically in resources (e.g., security groups).**

**Useful when the block count or content is variable.**

*Hint: Need to create N similar nested blocks with loops?*

## 7. “Explain the difference between depends\_on and implicit dependencies in Terraform.”

✓ Answer:

Implicit dependencies come from references; explicit ones are declared using depends\_on.

Use depends\_on when Terraform can't infer order.

Hint: *Sometimes Terraform doesn't know what should come first — help it out.*

## 8. “How do you migrate Terraform state between backends?”

✓ Answer:

Use `terraform init -migrate-state` to move state to a new backend.  
Ensure backups and access before migration.

Hint: *You're moving from local to S3 — how to avoid data loss?*

## 9. “What is the purpose of terraform import?”

✓ Answer:

`terraform import` maps existing resources to Terraform code without recreating them.

It helps bring unmanaged resources under control.

Hint: *Your infra already exists — now bring it under Terraform.*

## 10. “How do you handle breaking changes when upgrading Terraform provider versions?”

✓ Answer:

Read changelogs, use version constraints in required\_providers, and test upgrades in isolated environments.

Use terraform state replace-provider if needed.

Hint: *You upgraded the AWS provider and nothing works — now what?*

---

## SCENARIO-BASED INTERVIEW QUESTIONS

---

1. Ask: “*Your team reported that a recent Terraform apply caused service downtime. How do you prevent this in future runs?*”

✓ Expected: *Use lifecycle blocks like create\_before\_destroy, set depends\_on, and use terraform plan to review changes before apply.*

Hint: *What lifecycle setting ensures zero downtime when replacing resources?*

2. Ask: “*How do you manage Terraform state in a team with multiple contributors working on different modules?*”

✓ Expected: *Split state using separate backends per module or use workspaces. Use remote backend with locking (e.g., S3 + DynamoDB).*

Hint: *Think: isolation, locking, and concurrent safety.*

3. Ask: “*You have multiple modules with dependencies. Terraform is trying to destroy a dependent resource first. How do you fix it?*”

- ✓ Expected: *Use depends\_on to explicitly control the destroy/create order.*

Hint: *Implicit dependencies not enough? Time to be explicit.*

**4. Ask: “You’re upgrading a provider version and the plan shows unexpected changes. How do you handle provider upgrades safely?”**

- ✓ Expected: *Read changelogs, pin versions with required\_providers, and test in isolated branches or workspaces before merging.*

Hint: *Upgrades ≠ safe by default. How do you stay protected?*

**5. Ask: “You’ve been asked to detect manual infra changes made outside of Terraform. What’s your strategy?”**

- ✓ Expected: *Use terraform plan to detect drift or use tools like terraform plan -detailed-exitcode in CI/CD.*

Hint: *Terraform can’t fix what it doesn’t track. Plan is your detective.*

---

## PROJECT-BASED INTERVIEW QUESTIONS

---

**1. Ask: “How would you design a Terraform solution to deploy infrastructure across multiple AWS accounts using a single repo?”**

- ✓ Expected: *Use modules with account-specific variables, assume*

*roles for each account, and structure directories per environment/account.*

Hint: *Think: cross-account auth + reusability + separation.*

**2. Ask: “You need to provision Kubernetes clusters across GCP and AWS using Terraform. How would you design this setup?”**

❖ Expected: *Use separate providers and workspaces per cloud, with reusable modules per cluster type.*

Hint: *Multi-cloud needs clean separation and provider switching.*

**3. Ask: “Suppose you want to automate destroying old test environments every 7 days. How would you build this with Terraform?”**

❖ Expected: *Use a CI/CD pipeline with a script that checks resource age (via tags) and triggers terraform destroy for old workspaces.*

Hint: *Tag it, script it, automate it.*

**4. Ask: “Your Terraform plan applies 100+ resources and takes too long. How would you optimize execution time?”**

❖ Expected: *Use parallelism, break infra into smaller modules with isolated states, and avoid unnecessary refreshes.*

Hint: *Smaller chunks = faster runs. Don't refresh what's untouched.*